

Computer Vision 1

The background is a dark blue grid with various geometric elements. On the right side, there are several concentric circles and arcs, some solid and some dashed, in a lighter blue color. On the left side, there are horizontal lines with arrowheads pointing right, and some small white dots. The overall aesthetic is technical and futuristic.

THEO GEVERS

MASTER AI

UNIVERSITY OF AMSTERDAM

Lectures/Theory

- 06-02-2018, 17:00-19:00, C0.05, **Introduction** (*Szeliski 1*)
- 13-02-2018, 17:00-19:00, C0.05, **Image Formation** (*Szeliski: 2.1.1 + 2.1.2 + 2.2 + 2.3.2 + 2.3.3*)
- 20-02-2018, 17:00-19:00, C0.05, **Color and Image Processing** (*Szeliski: 3.1 + 3.2 + 3.3*)
- 27-02-2018, 17:00-19:00, C0.05, **Feature Detection, Motion and Classification** (*Szeliski: 4, 8.1.1 + 8.1.3 + 8.2.1 + 8.4; Bengio: 4 + 5.1 + 5.2 + 5.3 + 5.7 + 5.8 + 5.9*)
- 06-03-2018, 17:00-19:00, C0.05, **Object Recognition: BoW and Deep Learning** (*Szeliski: 5.1.1 + 5.1.4 + 5.1.5 + 5.2 + 5.3 + 5.4, 6.1 + 6.3, 14.1 + 14.2.1 + 14.3 + 14.4.1; Bengio: 7.2 + 7.4 + 9.1 + 9.2 + 9.3*)
- 13-03-2018, 17:00-19:00, C0.05, **ConvNets, Stereo and 3D Reconstruction** (*Szeliski: 11.1 + 11.2 + 11.3 + 11.4, 12.1 + 12.2; Bengio: 12.1 + 12.2*)
- 20-03-2018, 17:00-19:00, C0.05, **Applications** (*Szeliski: 12.6.2 + 12.6.3 + 12.2.4*)
- 26-03-2018, Monday, 9:00-12:00, **Written Exam**

Today's class



**Object Recognition
(ConvNets)**

Object Detection

Stereo Vision

Object Recognition using Deep Learning

Artificial Model (1943 McCulloch/Pitts)



The Neuron


Inputs: $\vec{x} = x_1, x_1, \dots, x_n$

Weights: $\vec{w} = w_1, w_1, \dots, w_n$


Logit:
$$z = \sum_{i=1}^n w_i x_i + b$$

Output: $y = f(z)$

Output: $y = f(\vec{x} \cdot \vec{w} + b)$



Linear Perceptron



**Example of a linear decision boundary
for binary classification.**


Inputs x_1, x_1, \dots, x_n

Weights w_1, w_1, \dots, w_n

Logit $z = \sum_{i=1}^n w_i x_i + b$

$$f(z) = \begin{cases} -1 & \text{if } z < 0 \\ 1 & \text{if } z \geq 0 \end{cases}$$

Linear Perceptron




Inputs x_1, x_1, \dots, x_n

Weights w_1, w_1, \dots, w_n

Logit $z = \sum_{i=1}^n w_i x_i + b$


$$f(z) = \begin{cases} -1 & \text{if } z < 0 \\ 1 & \text{if } z \geq 0 \end{cases}$$




Sigmoid, Tanh, and ReLU Neurons

Sigmoid


$$\sigma(x) = 1 / (1 + e^{-z})$$



tanh $\tanh(z)$



ReLU $\max(0, z)$



Feed-forward Neural Networks


$$\vec{y} = f(\vec{W} \cdot \vec{x} + \vec{b})$$

Weight matrix $n \times m$: \vec{W}

Inputs: $\vec{x} = x_1, x_1, \dots, x_n$

Outputs: $\vec{y} = y_1, y_1, \dots, y_m$

Bias: \vec{b}



Softmax Output Layer


$$y_i = \frac{e^{z_i}}{\sum_j e^{z_j}}$$

A mostly complete chart of Neural Networks


©2016 Fjodor van Veen - asimovinstitute.org

-  Backfed Input Cell
-  Input Cell
-  Noisy Input Cell
-  Hidden Cell
-  Probabilistic Hidden Cell
-  Spiking Hidden Cell
-  Output Cell
-  Match Input Output Cell
-  Recurrent Cell
-  Memory Cell
-  Different Memory Cell
-  Kernel
-  Convolution or Pool


Perceptron (P)



Feed Forward (FF)




Radial Basis Network (RBF)




Deep Feed Forward (DFF)




Recurrent Neural Network (RNN)




Long / Short Term Memory (LSTM)




Gated Recurrent Unit (GRU)




Auto Encoder (AE)




Variational AE (VAE)



Denoising AE (DAE)



Sparse AE (SAE)



Single layer feed-forward NN




Figure: Single layer feed forward NN

CIFAR-10

10 labels

50,000 training images, each image is tiny: 32x32

10,000 test images.

airplane



automobile



bird



cat



deer



dog



frog



horse



ship




truck





$$\vec{y} = f(\vec{W} \cdot \vec{x} + \vec{b})$$



10 numbers,
indicating class
scores

[32x32x3]

array of numbers

i.e. 3072 numbers total


Output=10x1

Weights=10x3072

Bias=10x1

$$\vec{y} = f(\vec{W} \cdot \vec{x} + \vec{b})$$

Input=3072x1



10 numbers,
indicating class
scores

[32x32x3]

array of numbers 0...1
(3072 numbers total)

Output=3x1 **Weights=3x4** **Bias=4x1**

$$\vec{y} = f(\vec{W} \cdot \vec{x} + \vec{b})$$

Input=4x1

(cat/dog/ship)

stretch pixels into single column



input image

0.2	-0.5	0.1	2.0
1.5	1.3	2.1	0.0
0	0.25	0.2	-0.3

\vec{W}

Weights=3x4

56
231
24
2

\vec{x}

Input=4x1

+

1.1
3.2
-1.2

\vec{b}

Bias=4x1

→

-96.8
437.9
61.95

\vec{y}

Output=3x1

cat score

dog score


ship score

Output=3x1 **Weights=3x4** **Bias=4x1**


$$\vec{y} = f(\vec{W} \cdot \vec{x} + \vec{b})$$

Input=4x1

(cat/dog/ship)




matrix multiply + bias offset



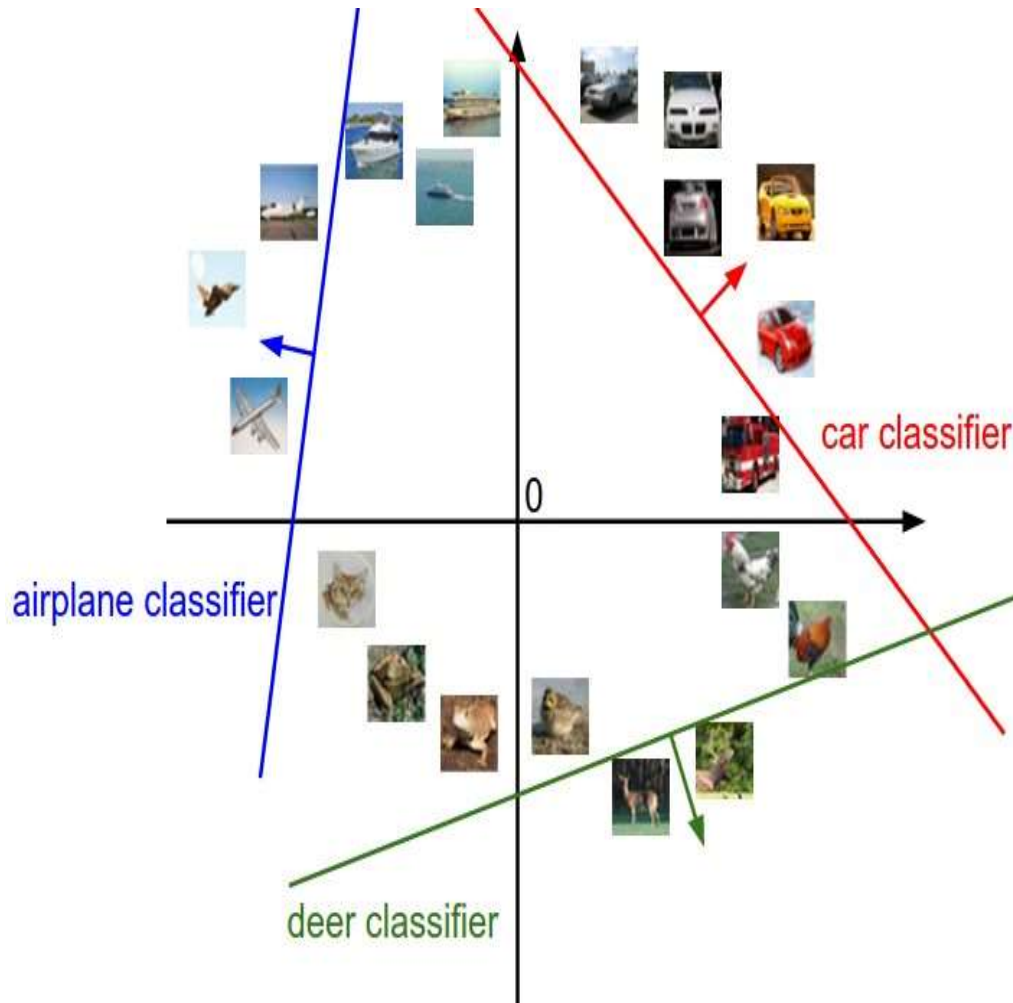
Softmax Output Layer

$$y_i = \frac{e^{z_i}}{\sum_j e^{z_j}}$$



\vec{y}

$$\vec{y} = f(\vec{W} \cdot \vec{x} + \vec{b})$$






[32x32x3]
array of numbers 0...1
(3072 numbers total)



$$f(x_i, W, b) = Wx_i + b$$

Example trained weights
of a linear classifier
trained on CIFAR-10:






"2-layer Neural Net", or
"1-hidden-layer Neural Net"


"3-layer Neural Net", or
"2-hidden-layer Neural Net"

"Fully-connected" layers


3 hidden neurons



6 hidden neurons



20 hidden neurons





more neurons = more capacity


Object Recognition using ConvNets

Convolutional Neural Networks (CNN, ConvNet, DCN)


- Fully connected NN's: e.g. input layer is $200 \times 200 \times 3$ color image = 120,000 weights!!
- CNN = a multi-layer neural network with
 - **Local** connectivity
 - **Share** weight parameters across spatial positions



Convolutional Layer (1)



Convolutional Layer (1)



Input volume of convolutional layer:

width w_{in} , height h_{in} , depth d_{in} , zero padding p


This input volume is processed by k filters. Filters are defined by:

spatial extent e , stride s

Output volume parameters:

width $w_{out} = \left(\frac{w_{in} - e + 2p}{s} \right) + 1$, height $h_{out} = \left(\frac{h_{in} - e + 2p}{s} \right) + 1$, depth $d_{out} = k$

Max Pooling




Pooling layer parameters:
spatial extent e , stride s

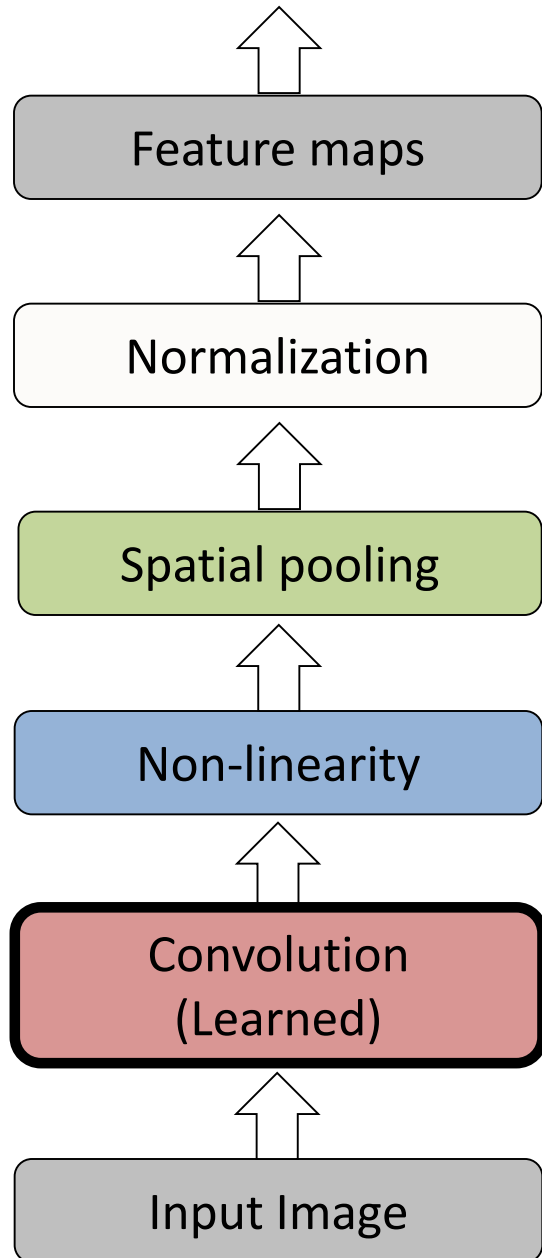
Resulting dimensions:

width $w_{out} = \left(\frac{w_{in} - e}{s} \right) + 1$, height $h_{out} = \left(\frac{h_{in} - e}{s} \right) + 1$

Convolutional Neural Networks



Convolutional Neural Networks

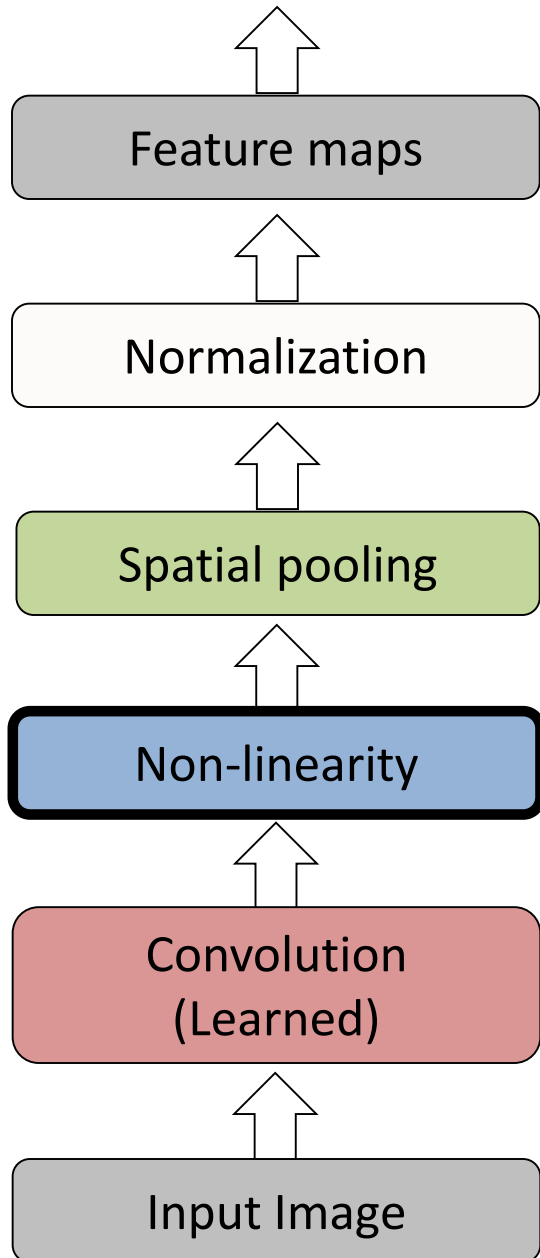


Input




Feature Map


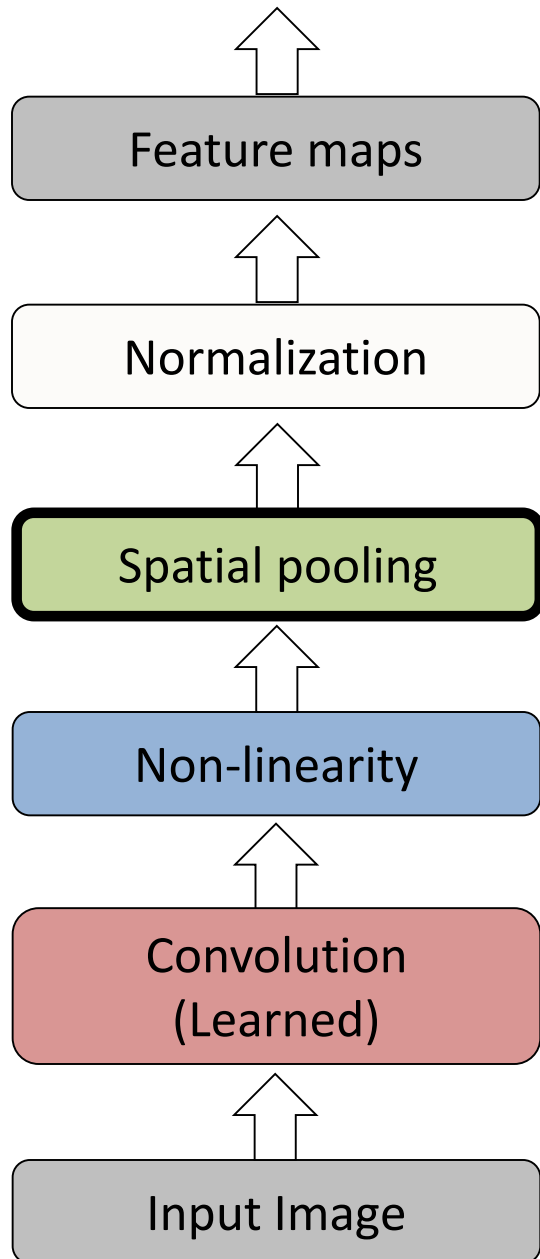
Convolutional Neural Networks



Rectified Linear Unit (ReLU)




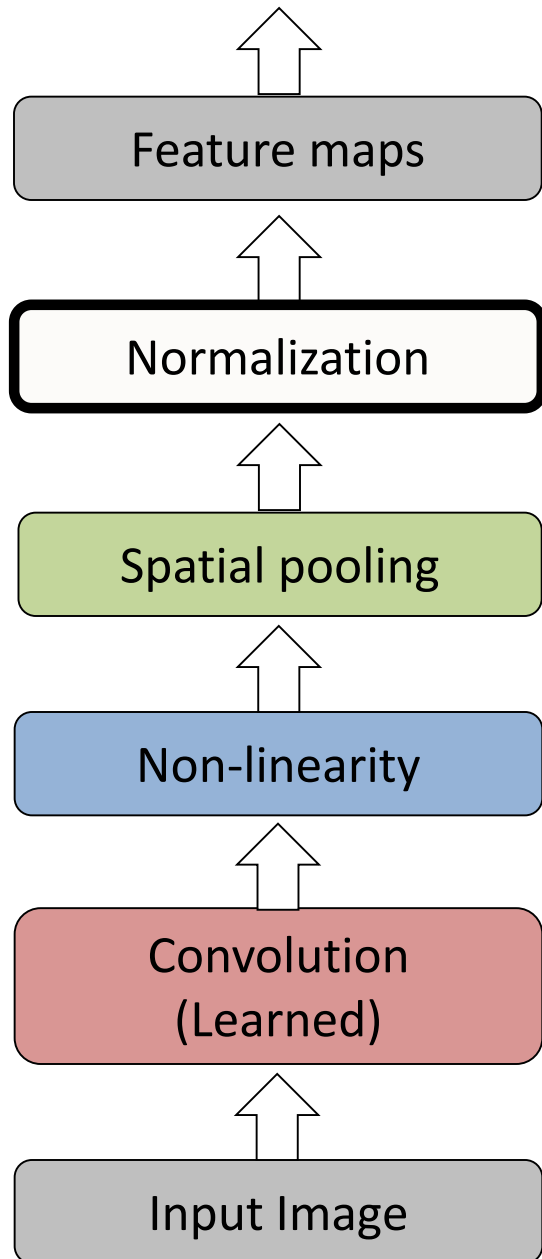
Convolutional Neural Networks




Max-pooling: a non-linear down-sampling

Provide *local invariance*

Convolutional Neural Networks

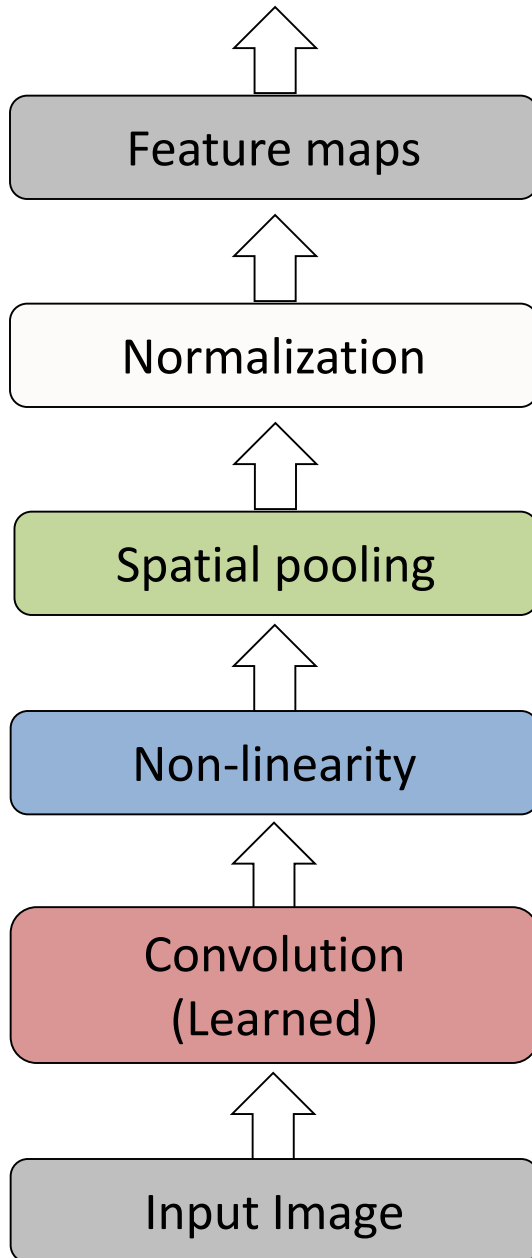


Feature Maps

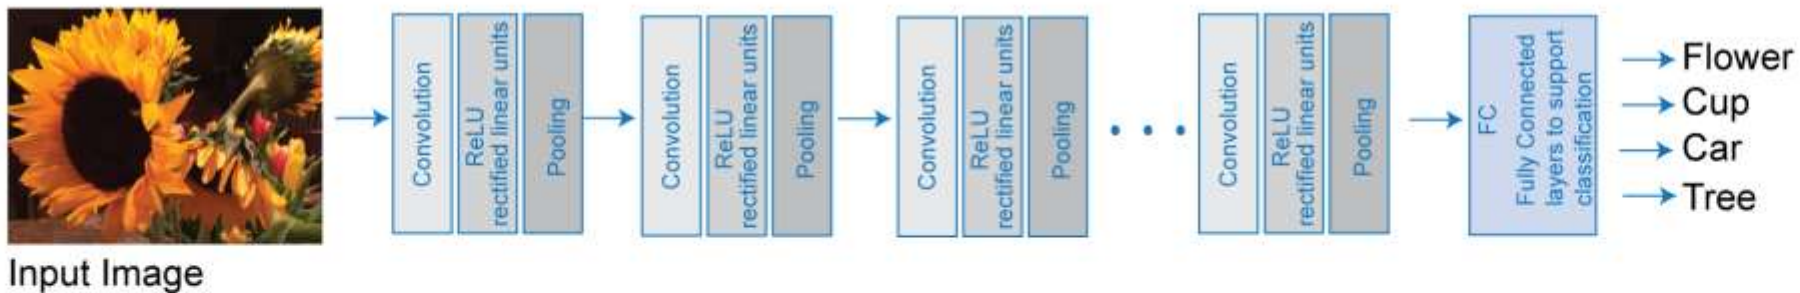


Feature Maps
After Contrast
Normalization


Convolutional Neural Networks



Standard ConvNets Architecture




VGGNet (2014)



VGGNet 16

LeNet [LeCun et al. 1998]




Gradient-based learning applied to document recognition [[LeCun, Bottou, Bengio, Haffner 1998](#)]



LeNet-1 from 1993

AlexNet (2013 ImageNet Winner)

- Similar framework to LeCun'98 but:
 - Bigger model (8 hidden layers, 650,000 units, 60,000,000 params)
 - More data (10^6 vs. 10^3 images)
 - GPU implementation (50x speedup over CPU)
 - Trained on two GPUs for a week

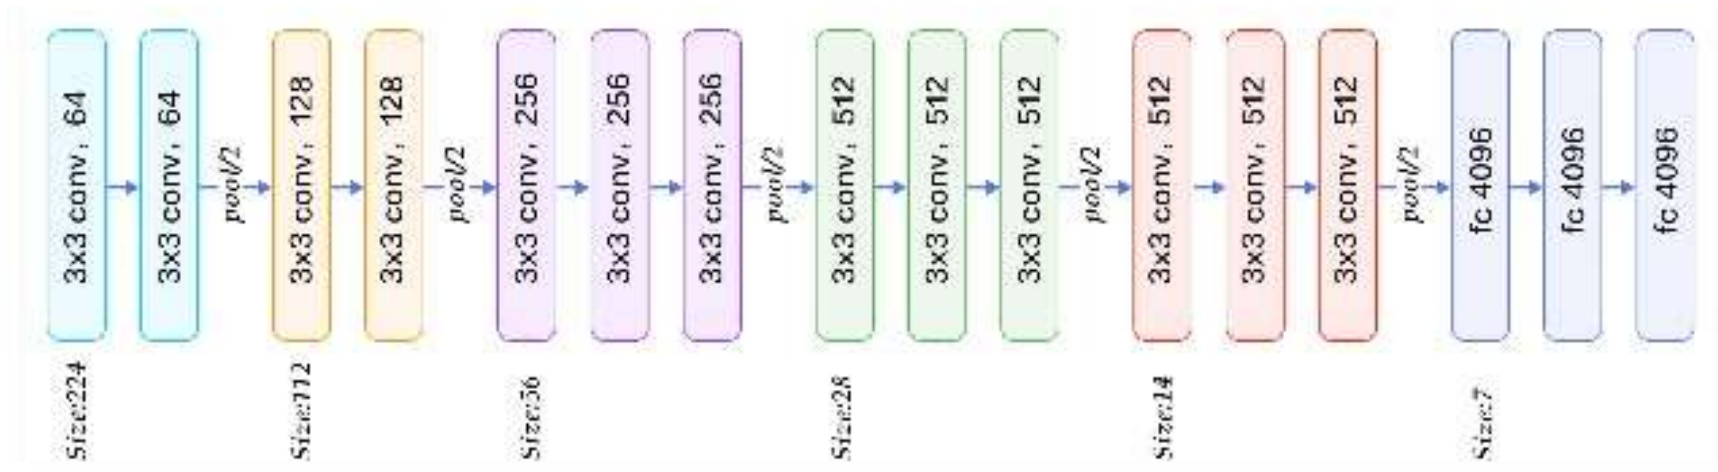


A. Krizhevsky, I. Sutskever, and G. Hinton,

[ImageNet Classification with Deep Convolutional Neural Networks](#), NIPS 2012

VGG Net (2014 ImageNet Winner)

VGGNet (2014)



VGGNet 16

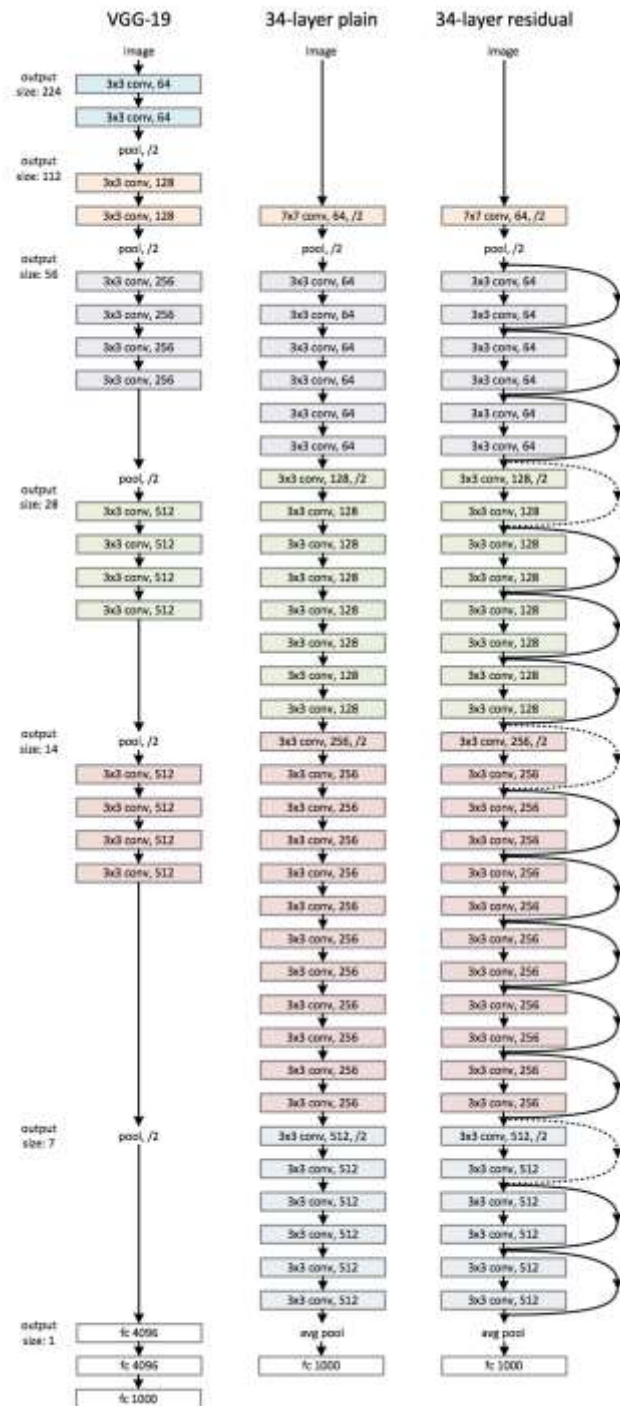
AlexNet and VGGNet

AlexNet





VGG16



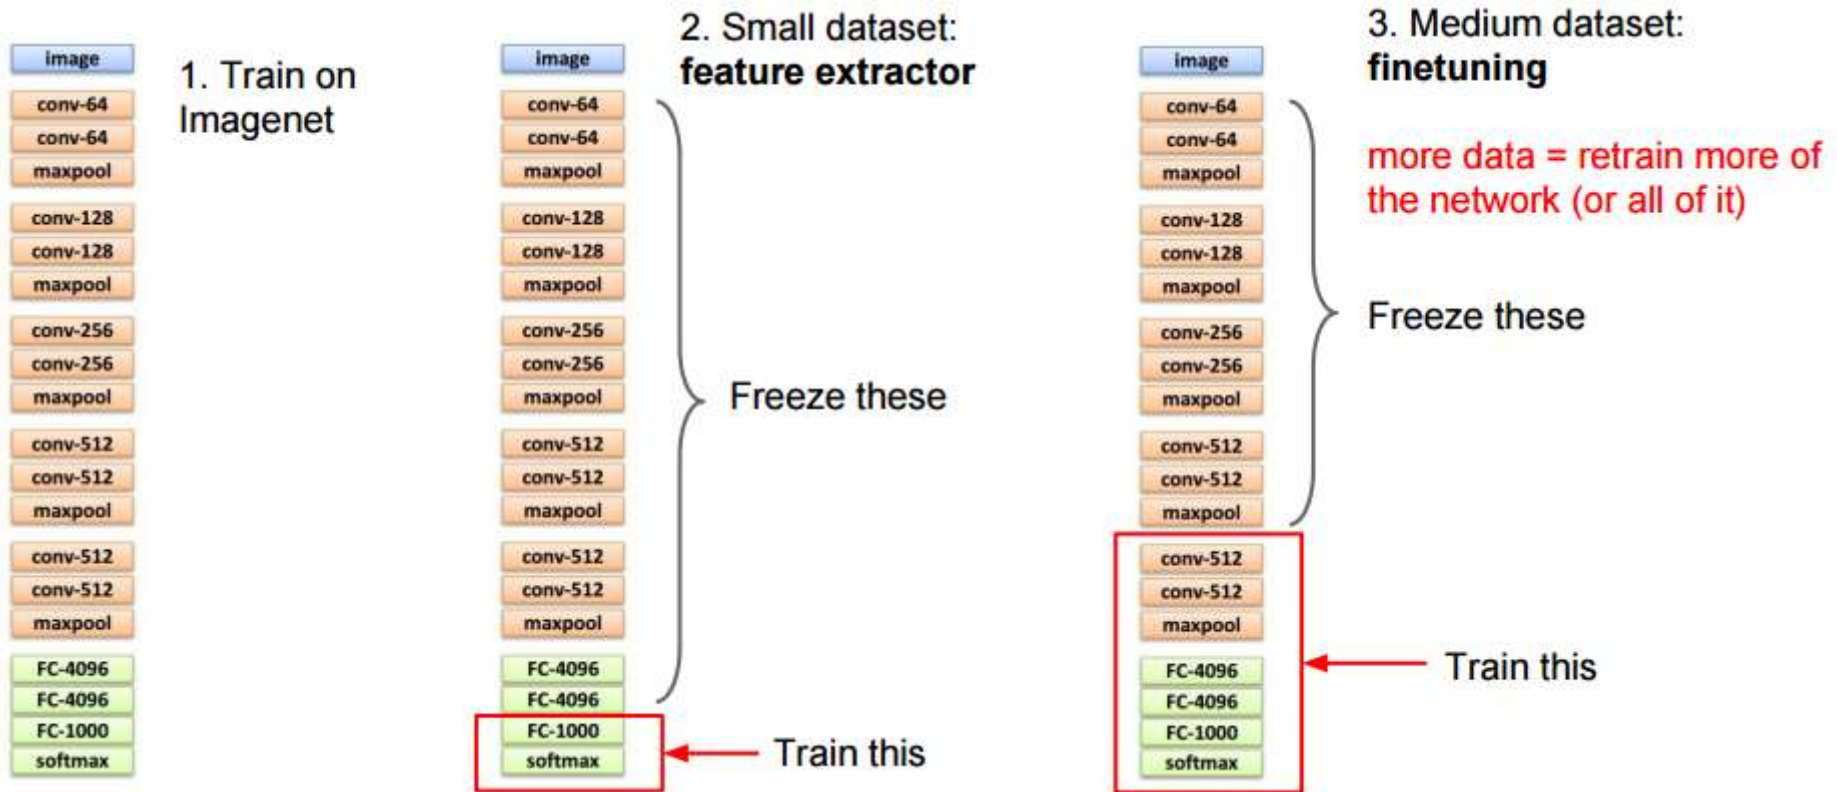


Features



Transfer Learning

- Improvement of learning in a **new** task through the *transfer of knowledge* from a **related** task that has already been learned.




Training Convolutional Neural Networks


- Backpropagation + stochastic gradient descent with momentum
 - [Neural Networks: Tricks of the Trade](#)
- Data augmentation
- Dropout
- Batch normalization
- Initialization
 - Transfer learning

Data Augmentation (Jittering)

- Create *virtual* training samples
 - Horizontal flip
 - Random crop
 - Color casting
 - Geometric distortion




ImageNet Challenge 2012-2016



ImageNet Classification

- 1000 categories and 1.2 million training images




Li Fei-Fei: ImageNet Large Scale Visual Recognition Challenge, 2014 <http://image-net.org/>

8/1/2014

lipa.p@nyu.edu


7

ILSVRC top-5 error on ImageNet



E2E: Classification: ResNet

Revolution of Depth



He, Kaiming, Xiangyu Zhang, Shaoqing Ren, and Jian Sun. "Deep Residual Learning for Image Recognition." arXiv preprint arXiv:1512.03386 (2015). <https://arxiv.org/abs/1512.03386>

Tools

- [Caffe](#)
- [cuda-convnet2](#)
- [Torch](#)
- [MatConvNet](#)
- [Pylearn2](#)
- [TensorFlow](#)

Today's class


Object Detection

Stereo

3D Reconstruction

Window-based models

Generating and scoring candidates



Discriminative Classifier Construction


Nearest neighbor



10^6 examples

Shakhnarovich, Viola, Darrell 2003
Berg, Berg, Malik 2005...


Neural networks



AlexNet 2013
VGGNet 2014


...

Support Vector Machines




Guyon, Vapnik
Heisele, Serre, Poggio,
2001,...

Boosting



Viola, Jones 2001,
Torralba et al. 2004,
Opelt et al. 2006,...

Conditional Random Fields



McCallum, Freitag, Pereira
2000; Kumar, Hebert 2003
...

The Viola/Jones Face Detector

- A seminal approach to real-time object detection
- Training is slow, but detection is very fast
- Key ideas
 - *Integral images* for fast feature evaluation
 - *Boosting* for feature selection
 - *Attentional cascade* for fast rejection of non-face windows

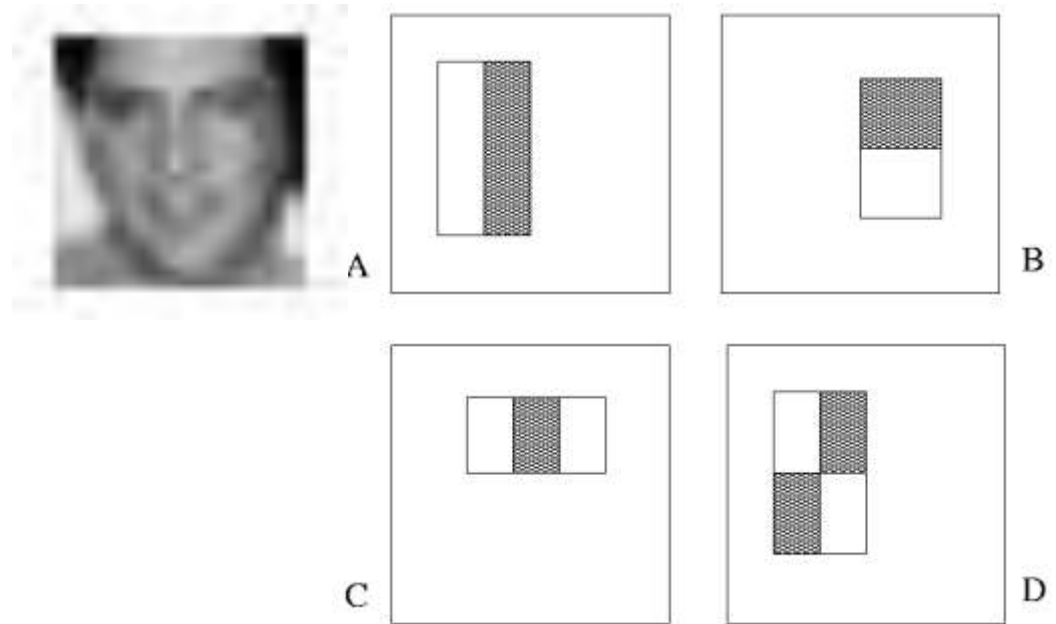
P. Viola and M. Jones. [Rapid object detection using a boosted cascade of simple features](#). CVPR 2001.

P. Viola and M. Jones. [Robust real-time face detection](#). IJCV 57(2), 2004.

~14000 citations!

Image Features

“Rectangle filters”




Value =

$$\sum (\text{pixels in white area}) - \sum (\text{pixels in black area})$$

Fast computation with integral images

- The *integral image* computes a value at each pixel (x,y) that is the sum of the pixel values above and to the left of (x,y) , inclusive
- This can quickly be computed in one pass through the image




Computing Sum within a Rectangle





- Let A,B,C,D be the values of the integral image at the corners of a rectangle
- Then the sum of original image values within the rectangle can be computed as:

$$\text{sum} = A - B - C + D$$


- Only 3 additions are required for any size of rectangle!



Output of Face Detector on Images



Profile Detection




Summary: Viola/Jones Detector

- Rectangle features
- Integral images for fast computation
- Boosting for feature selection
- Attentional cascade for fast rejection of negative windows

Window-based Detection: Strengths

- Sliding window detection and global appearance descriptors:
 - Simple detection protocol to implement
 - Good feature choices critical
 - Past successes for certain classes




Window-based Detection: Limitations

- High computational complexity
 - For example: 250,000 locations x 30 orientations x 4 scales = 30,000,000 evaluations!
 - If training binary detectors independently, means cost increases linearly with number of classes
- With so many windows, false positive rate better be low

Limitations (continued)

- Not all objects are “box” shaped



Segmentation as Selective Search for Object Detection

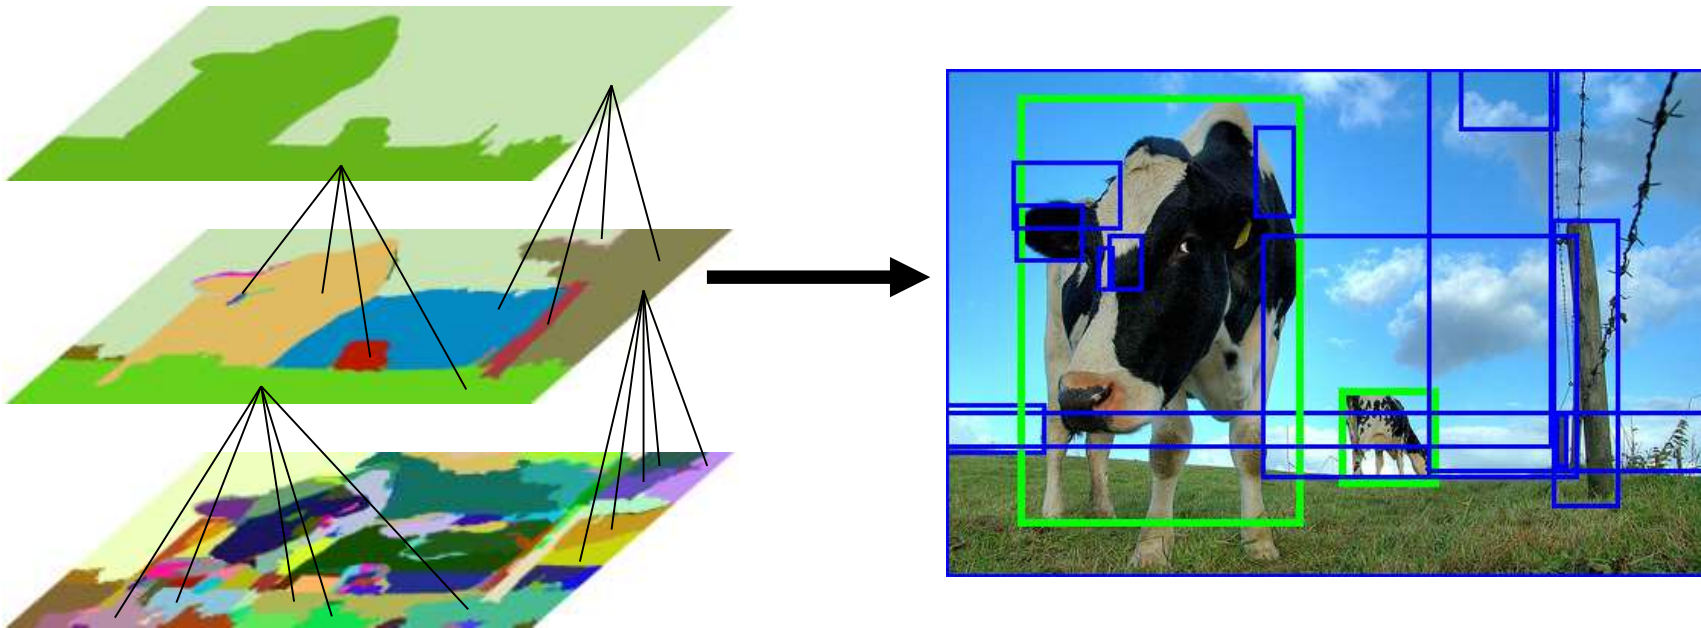
Jasper Uijlings, Koen van de Sande, Theo Gevers, Arnold Smeulders:
Selective Search for Object Recognition. International Journal of
Computer Vision 104(2): 154-171 (2013)

Selective Search for Recognition

- Design criteria
 - High recall
 - Coarse locations are sufficient
 - ⇒ Bounding boxes
 - Fast to compute
 - ⇒ Efficient low-level features
 - ⇒ <10s per image


Selective Search: Approach

- Hypotheses based on hierarchical grouping








Group adjacent regions on color/texture cues

Example 1



Example 2



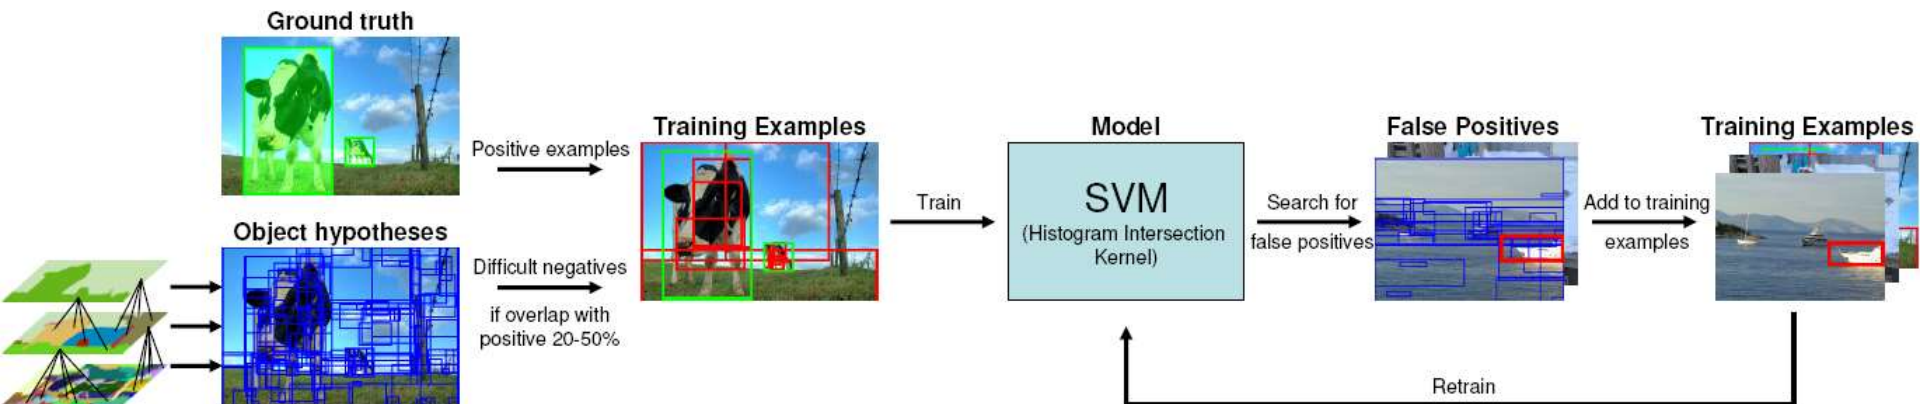
Selective Search on ILSVRC2011

- Apply to ILSVRC2011 train set
- Object hypotheses are class-independent

	ILSVRC2011 train
With bounding box annotations	315,525 images
Average #boxes/image	1,565
Average recall	98.5%

Localisation System Training

- Use positives and mirrored positives
- Use object hypotheses to create difficult initial negatives (at most 7,500)
- Add 2 iterations of false positives (from 4,000 images)



- Features: Bag-of-words, sample every pixel, SIFT, “ColorSIFT” and RGB-SIFT, pyramid up to level 3, codebook size 4096
- Histogram Intersection Kernel with Fast Approximation

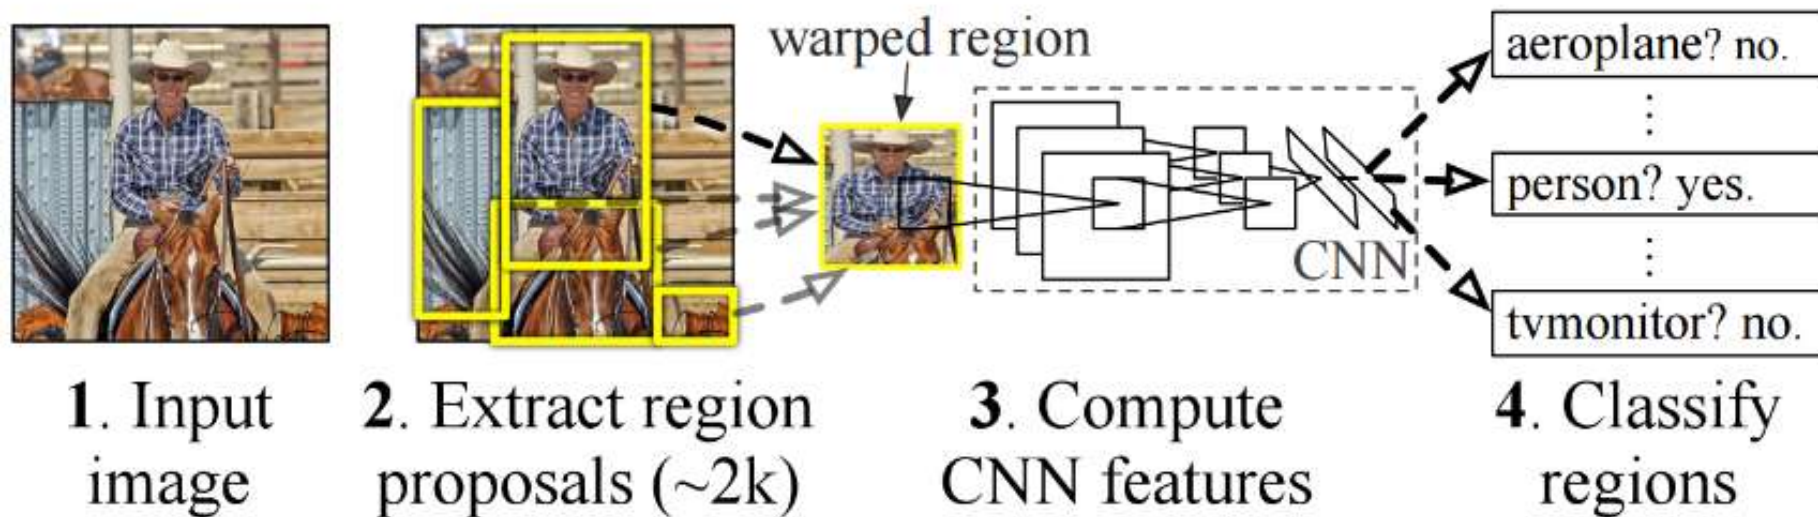
Summary

- Adopted segmentation as selective search strategy for object localisation:
 - High recall: >96% with ~1,500 locations
 - Coarse locations are sufficient: bounding boxes
 - Fast to compute: <10s per image
 - Class-independent
 - Enables the use of bag-of-words features


R-CNN: Regions with CNN features

- Trained on ImageNet classification
- Finetune CNN on PASCAL

R-CNN: *Regions with CNN features*



Fast R-CNN



Fast RCNN [[Girshick, R 2015](https://arxiv.org/abs/1504.08055)]
<https://github.com/rbgirshick/fast-rcnn>

Faster R-CNN

1. Input image




2. Face detection




Mathias et al. detector

3. Cropped face




+ 40% margin

4. Feature extraction



VGG-16 architecture

5. Prediction



Softmax expected value

$$\Sigma = 23.4 \text{ years}$$



Figure 5. Sample detection results on the FDDB dataset, where green bounding boxes are ground-truth annotations and red bounding boxes

Application: Deep Learning Image Intrinsics

Recovering Lightness: Retinex

- Image Intensity: $I = e\rho$

- Take Logarithm: $\log I = \log e + \log \rho$

OR $I = e + \rho$

- Use Laplacian:

$$d = \nabla^2 I = \nabla^2 e + \nabla^2 \rho \quad \nabla^2 = \frac{\partial^2}{\partial x^2} + \frac{\partial^2}{\partial y^2}$$


- Sharp changes in reflectance ρ

$\nabla^2 \rho$ has 2 infinite spikes near edges and $\nabla^2 \rho = 0$ elsewhere

- Smooth changes in illumination e

$\nabla^2 e \approx 0$ everywhere

Solving the Inverse Problem



Find lightness $l(x, y)$ from $t(x, y)$:

Poisson's Equation

$$\nabla^2 l = t$$

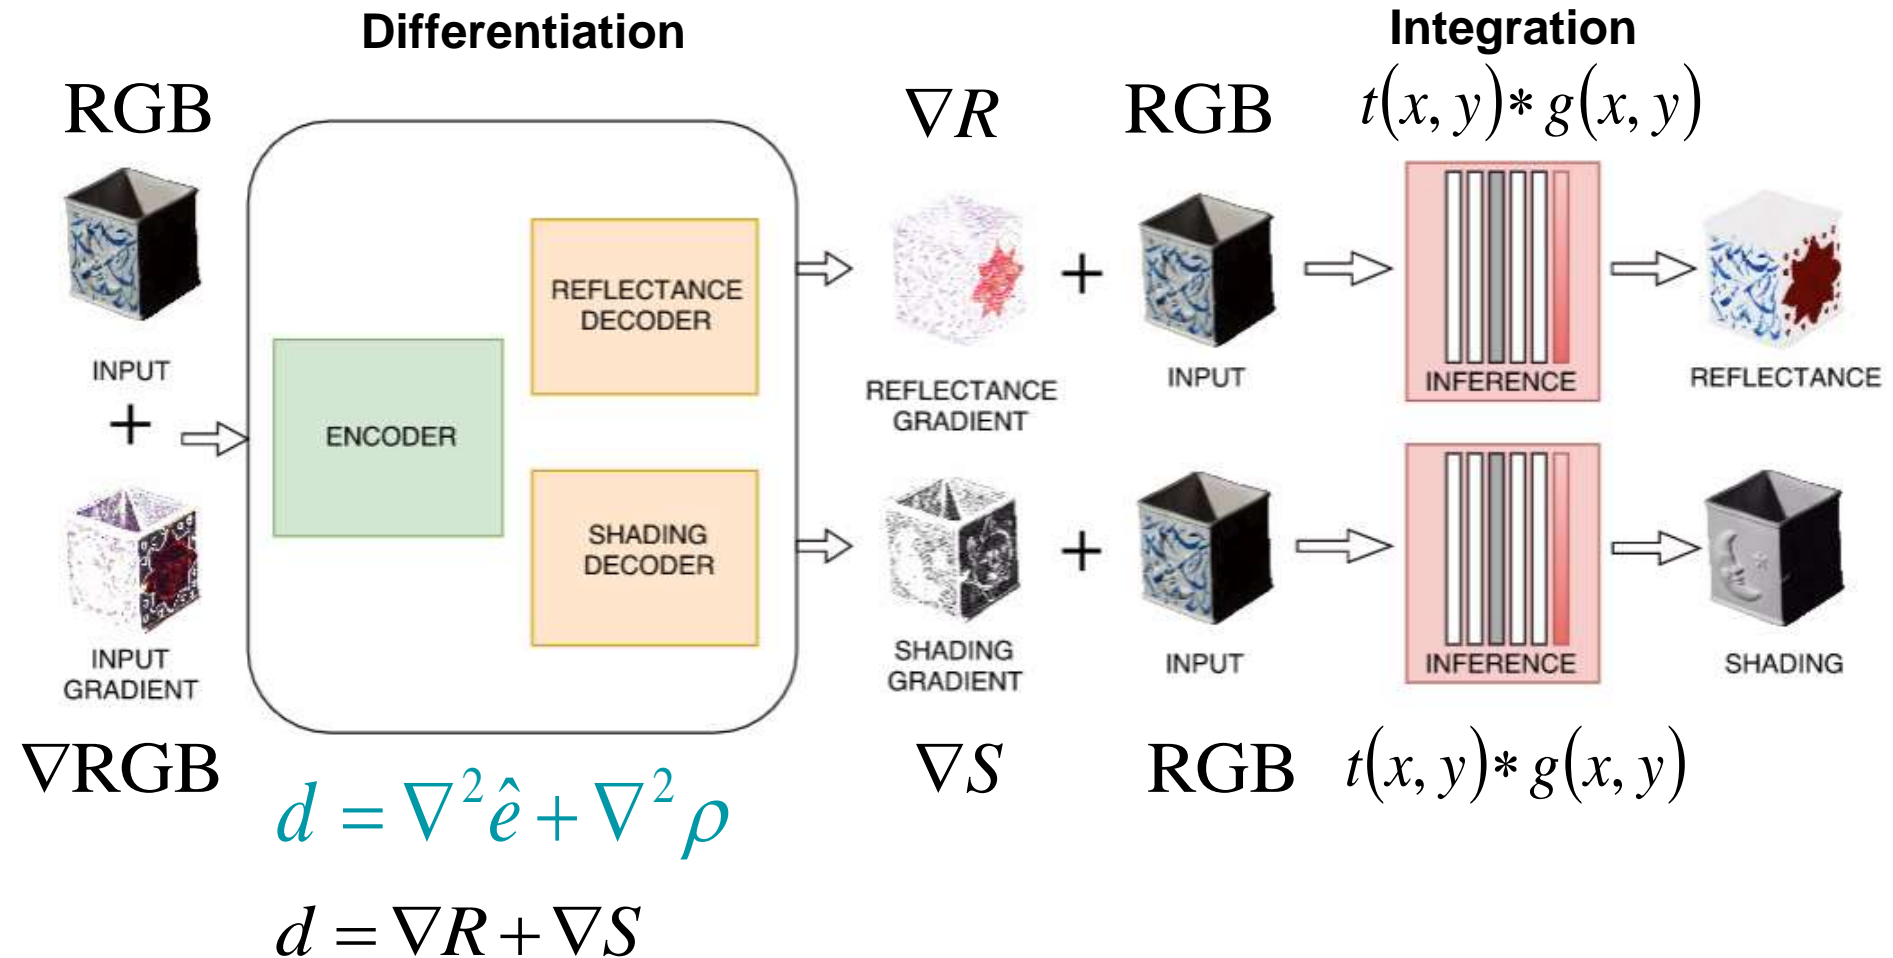
$$\left(\frac{\partial^2}{\partial x^2} + \frac{\partial^2}{\partial y^2} \right) l(x, y) = t(x, y)$$

We have to find $g(x, y)$ which satisfies



$$l(x, y) = \int_{-\infty}^{\infty} \int_{-\infty}^{\infty} t(u, v) g(x - u, y - v) du dv$$

$$l(x, y) = t(x, y) * g(x, y)$$

RetiNet: Retinex-Inspired ConvNet



Qualitative Results: MIT Dataset



RetiNet: Conclusion

Advantage

- + Can capture most of the color information & shading
- + Colors are more vivid and eliminates most of the color artifacts
- + Fast and memory efficient
- + Sharp Edges

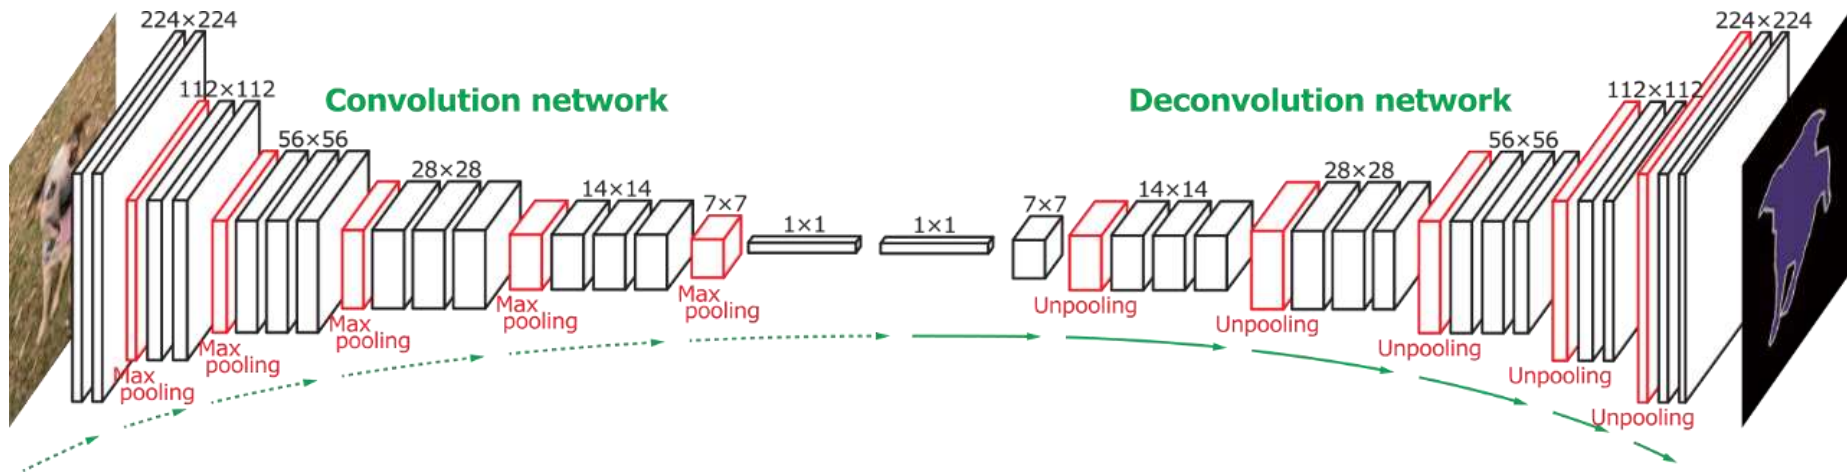
Application

- Semantic Image Segmentation

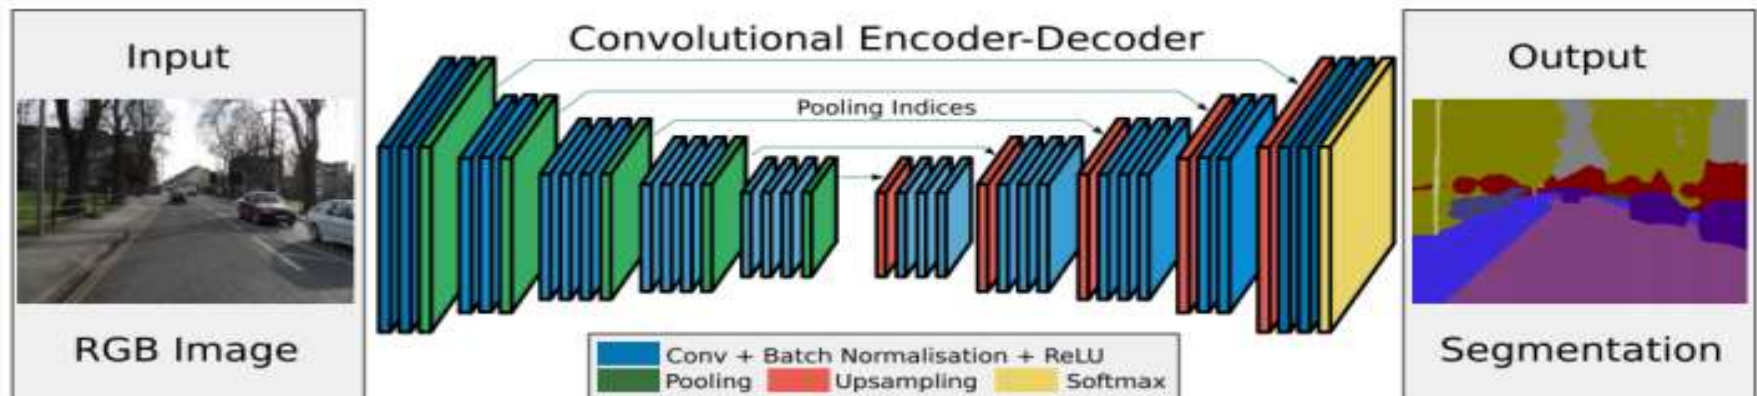
Application: Semantic Segmentation and Style Transfer

Semantic Segmentation

- Deep fully convolution network: SegNet
 - Encoder: learn low-level features in images
 - Decoder: reconstruct the image structure with labels



Semantic Segmentation: SegNet



Badrinarayanan, et al., "SegNet: A deep convolutional Encoder-Decoder Architecture for Image Segmentation."

Style Transfer



Application: Face Analysis

Emotions



Alamy BCC-RT

Facial Action Units

- Charles Darwin, 1872
- Ekman & Friesen, 1978



AU 1
Inner brow raise



AU 2
Outer brow raise



AU 4
Brow lower



AU 6
Cheek raise



AU 9
Nose wrinkler



AU 12
Lip corner pull



AU 15
Lip corner depress



AU 20
Lip stretcher



AU 23
Lip tight



AU 28
Lip suck




AU 37
Lip wipe




AU 84
Head shake
back and forth


Face Analysis



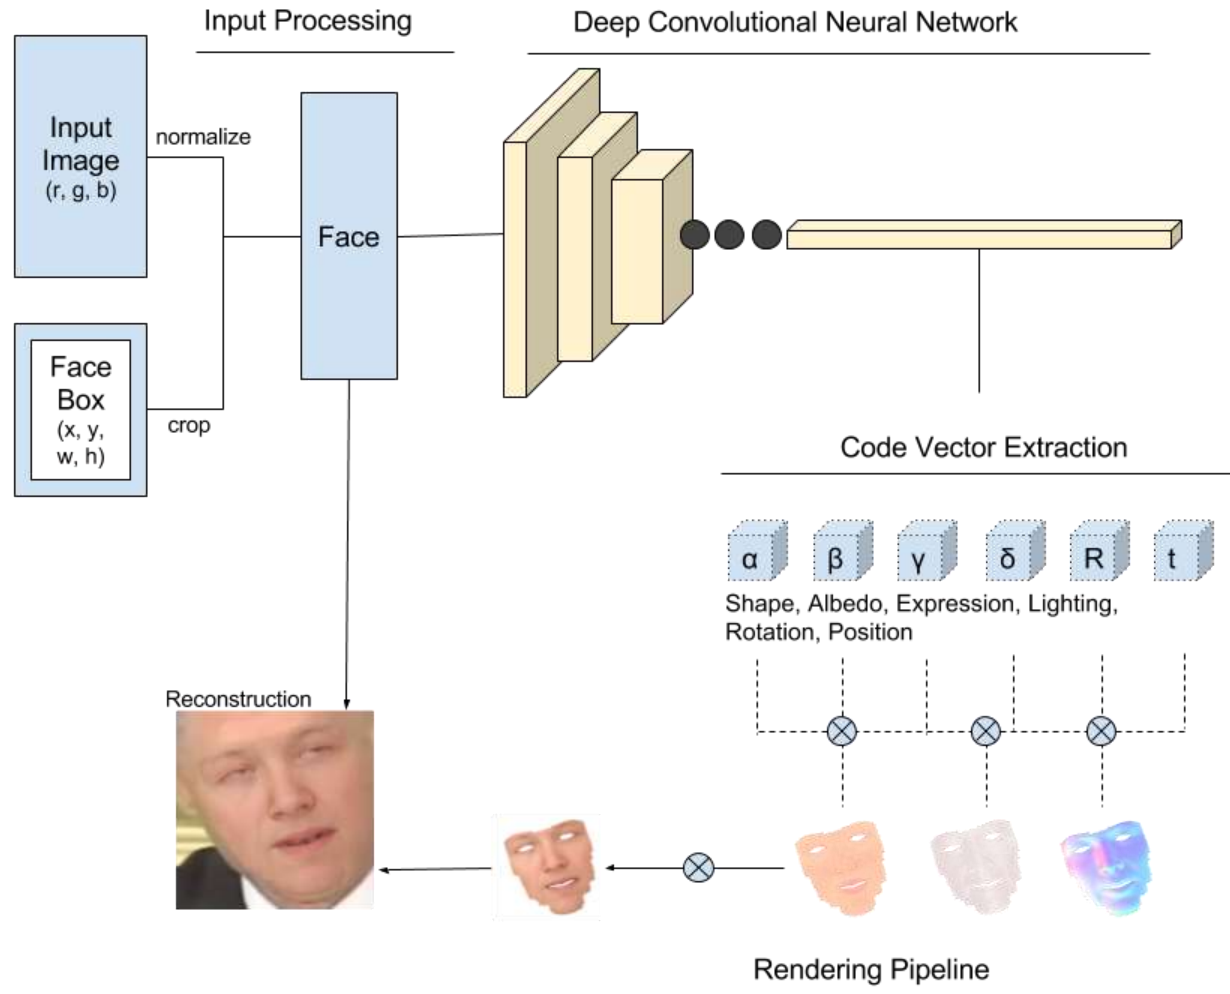
Face Analysis by Deep Learning


















Sightcorp



CNN Model: Multi-Task Supervised



CNN Model: Results

Input Image	Albedo	Lighting	Surface Normals	Reconstruction
				
				
				

Conclusion:

Deep Learning in Computer Vision

- Object classification, detection and segmentation
- Optical flow
- Color constancy
- Intrinsic image decomposition
- 3D and slam
- Human behavior analysis
- Etc...

Computer Vision 1

(total #slides 89 / Lecture 6)

Summary

- 1. Object Recognition (ConvNets)**
- 2. Object Detection**
- 3. Stereo Vision (Next Week)**